




HTTP

Junghoo Cho
cho@cs.ucla.edu

Web Interaction Example

- <http://www.amazon.com>

How Does It Work?



Core Internet Standards

- DNS (domain name service)
 - Internet protocol to map domain names to IPs
 - ICANN manages TLD (top-level domains)
- TCP/IP (transmission control protocol and internet protocol)
 - Internet routing and transportation protocol

Core Internet Standards

- HTTP (hypertext transportation protocol)
 - Communication protocol between web servers and web clients
- Encoding standards
 - Text: ASCII, Unicode (UTF-8)
 - Multimedia: JPEG, MP3, H.264, ...
- MIME (multipurpose internet mail extensions)
 - Data-type-specification standard

Core Internet Standards

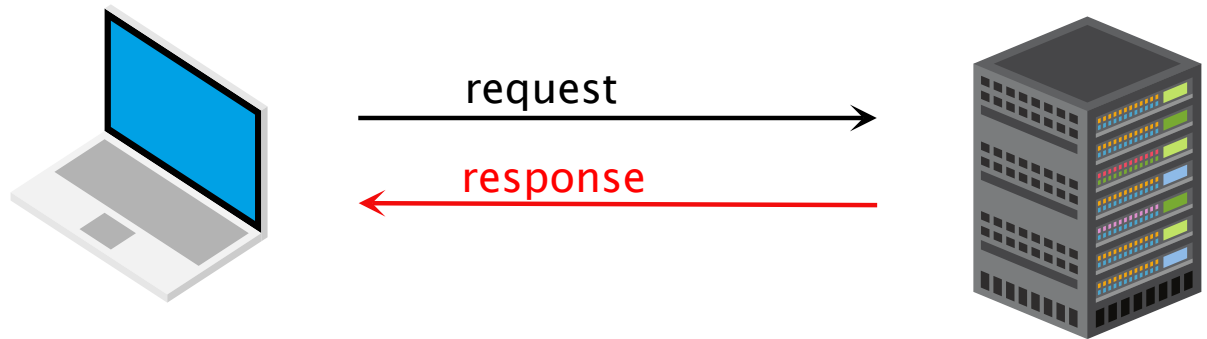
- HTML (hypertext markup language)
 - Markup standard
- CSS (Cascading style sheet)
 - Styling and formatting standard
- Javascript
 - De facto Web programming language

HTTP

- HTTP/2 (2015) is most recent
- HTTP/1.1 (1996) is still extremely popular
 - Major browsers support HTTP/2 only over HTTPS
- Learn HTTP/1.1 first
 - Then talk about HTTP/2
- Two key properties of HTTP
 - Request & response paradigm
 - Stateless protocol

HTTP: Request & Response

- All interactions start with a client's request



HTTP: Stateless Protocol

- Every request is handled independently of others
 - The server is not required to remember “history” of past requests
- Questions
 - Q: What are pros/cons of stateless protocol?
 - Q: what are the implications?

HTTP Example

- Telnet to <http://oak.cs.ucla.edu/classes/cs144>
- Real HTTP Request by browser

HTTP Message

- A message is either a request or response
- Message structure
 - request/status line
 - header
 - [empty line]
 - body
- Bare minimum HTTP request:

GET / HTTP/1.0

HTTP Request Example

```
GET /apps/echo.py HTTP/1.1
Host: oak.cs.ucla.edu
Connection: keep-alive
Cache-Control: max-age=0
Dnt: 1
Upgrade-Insecure-Requests: 1
User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10_15_3) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/80.0.3987.122 Safari/537.36
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,*/*;q=0.8
Accept-Encoding: gzip, deflate
Accept-Language: en-US,en;q=0.9,ko;q=0.8
Cookie: _xsrf=2|11b01aa1
```

HTTP Response Example

```
HTTP/1.1 200 OK
Date: Wed, 04 Apr 2019 03:20:33 GMT
Cache-Control: max-age=0, no-cache, s-maxage=10
Connection: Keep-Alive
Content-Encoding: gzip
Content-Length: 997
Content-Type: text/html; charset=UTF-8
Server: Apache/2.4.29 (Ubuntu)
Vary: Accept-Encoding
```

```
<html>
<head><title>Example page</title></head>
...
```

HTTP Request Line

- METHOD PATH PROTOCOL_VERSION

GET /apps/echo.py HTTP/1.1

HTTP Methods (1)

- **GET**: “retrieve” data
 - **IMPORTANT**: NO significant side effect at server
- **POST**: “post” data at the specified URL
 - May leave a side-effect on the server
- **PUT**: “place” the data at the URL (~ replace the data)
- **DELETE**: “delete” the data at the URL

HTTP Methods (2)

- Less common:
 - **HEAD**: same as **GET** but requests header only
 - **OPTIONS**: requests information on options available at the server
 - **TRACE**: the final recipient returns the whole request message in the body

HTTP Header

```
GET /apps/echo.py HTTP/1.1
Host: oak.cs.ucla.edu
Connection: keep-alive
Cache-Control: max-age=0
Dnt: 1
Upgrade-Insecure-Requests: 1
User-Agent: Mozilla/5.0 (Macintosh; Intel Mac OS X 10_15_3) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/80.0.3987.163 Safari/537.36
Accept: text/html,application/xhtml+xml,application/xml;q=0.9,image/webp,*/*;q=0.8
Accept-Encoding: gzip, deflate
Accept-Language: en-US,en;q=0.9,ko;q=0.8
Cookie: _xsrf=2|11b01aa1
```

HTTP Response

```
HTTP/1.1 200 OK
Date: Wed, 04 Apr 2019 03:20:33 GMT
Cache-Control: max-age=0, no-cache, s-maxage=10
Connection: Keep-Alive
Content-Encoding: gzip
Content-Length: 997
Content-Type: text/html; charset=UTF-8
Server: Apache/2.4.29 (Ubuntu)
Vary: Accept-Encoding
```

```
<html>
<head><title>Example page</title></head>
...
```

HTTP Response

- Structure
 - Status line
 - Header
 - [empty line]
 - body
- Status line
 - `VERSION STATUS_CODE REASON_PHRASE`
 - Example: `HTTP/1.1 200 OK`

Status Code

- **2xx**: Success
- **3xx**: Redirection
- **4xx**: Client Error
- **5xx**: Server Error

HTTP Response Header

```
HTTP/1.1 200 OK
Date: Wed, 04 Apr 2019 03:20:33 GMT
Cache-Control: max-age=0, no-cache
Connection: Keep-Alive
Content-Encoding: gzip
Content-Length: 997
Content-Type: text/html; charset=UTF-8
Server: Apache/2.4.29 (Ubuntu)
Vary: Accept-Encoding

<html>
<head><title>Example page</title></head>
...
```

Questions?

HTTP/2: Background

- Many high-latency limited-bandwidth mobile devices
- Many objects are needed to display a single page
 - HTML, image, CSS, JavaScript, ...
 - ~ 100 objects, ~ 2MB

HTTP/2

- Key assumptions have been relaxed for speed and efficiency
 - Server push
 - Stateful header compression
 - Binary encoding
 - ...

HTTP/2 New Features (1)

- “Multiplexed streams”
 - Multiple outstanding requests through a single connection
 - Split messages into small frames
 - Priority specification
- HPACK
 - Stateful HTTP header compression

HTTP/2 New Features (2)

- “Server push”
 - Predictively cache pushed by the server
- Binary format
- More detail at <https://daniel.haxx.se/http2/>

References

- HTTP/1.1: [RFC 7230](#) – [RFC 7237](#)
- HTTP/2: [RFC 7540](#)

